Chih-Kai Wang

Software engineering & AI internship

f0909172434@gmail.com
github.com/f0909172434
Taipei, Taiwan

I build Python and TypeScript tools for inspectable AI and math research: a CI check that reads the test report, not the exit code; an audit that ties claims to the bytes of their evidence; a counterexample search that reports the range it searched. I develop with Claude Code and Codex, review every change, and keep negative results public.

2	3	1,669 / 1,669	4
packages on npm and PyPI	pull requests merged upstream	released inputs accepted	negative results kept public

Education

National Taipei University of Education — Expected 2028
B.S. student, Mathematics Division, Department of Mathematics and Information Education.

Projects

HonestCI TypeScript · Node.js · GitHub Actions — v1.0.6 · npm, GitHub Marketplace

- Checks that JUnit reports are present, fresh, and not below a trusted test-count baseline.
- In the launch demo, a zero-test report with exit code 0 passes a bare workflow; with HonestCI it fails as HCI004_ZERO_TESTS. It does not judge test quality.

RigorGraph Python · JSON Schema · HTML reports — 1.0.1 · PyPI, public beta

- Local-first CLI that links research claims to evidence files and independent review records, checks SHA-256 hashes, and writes offline audit reports.
- Changed evidence bytes raise RG_HASH_MISMATCH, and a changed claim invalidates the old acceptance. VERIFIED means accepted by the recorded workflow, not true.

Finite Witness JavaScript · Web Workers · WebMCP · Python — v1.3.0 · browser tool

- Exhaustively searches small graphs (up to 6 vertices) for counterexamples and writes certificates that an independent Python script replays.
- Default run: 39 candidates examined, then a four-cycle with no triangle. Surviving a bounded search is evidence, not proof.

SAIR Proof Press Python · Lean 4 — public companion · frozen artifacts

- Companion to an equational-implication solver: Lean-checked proofs or finite countermodels.
- Frozen artifacts accepted 1,669 / 1,669 released inputs with zero model calls. This is released-input coverage, not a private score or rank.

Research

- **RuleDiff negative result**: four-page technical report. A lexical predictor scored 0.99 macro-F1 on development data and 0.67 on a held-out packet. Not peer reviewed.
- **RuleShift**: research pilot in which simple retrieval matched more complex agent memory strategies.
- **ProofWeave Core**: checks author-written Markdown proofs against pinned Lean 4 / Mathlib and reports certification apart from human-confirmed alignment.

Merged Upstream

- **DeepSeek Harness Desktop #740**: worktree path normalisation — Sep 2026
- **dsh-engram #4**: conservative database migration — Sep 2026
- **Codex Dream Skin #10**: Windows verification failures kept distinguishable — Jul 2026

Creative Work

ORACLE (short film), The Disease Called AI (song and watercolour MV) and world.execute(me) (terminal PV), October 2026: each rendered entirely from code.

Skills

Languages	Python, TypeScript, Node.js, Lean 4 (learning), Bash
Tools	GitHub Actions, pytest, Vitest, JSON Schema, MLX, Blender, Claude Code, Codex